



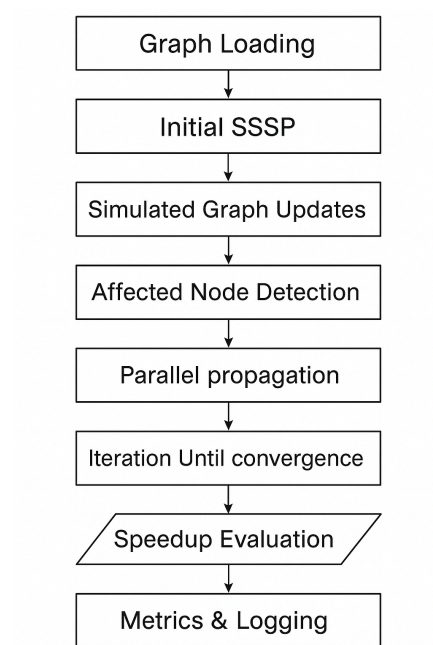
Member 1: Mustafa Irfan (i210626)

Member 2: Walia Fatima (i210838)

Member 3: Hasssan Qadir (i210883)

Section: G

## OpenMP – Based Dynamic SSSP Implementation



Our implementation follows the algorithmic principles outlined in the research paper:

**"A Parallel Algorithm Template for Updating Single-Source Shortest Paths in Large-Scale Dynamic Networks"**

The main idea is to avoid recomputing all shortest paths from scratch after a change in the graph (such as an edge insertion or deletion). Instead, this implementation:

- 1. **Runs Dijkstra's algorithm once at the start.**
- 2. **Detects affected nodes** when edges are inserted or deleted.
- 3. **Propagates changes to affected regions** using OpenMP parallelism.

**Step-by-Step Logic Flow**

Phase	Description	Relation to Paper
1.Graph Loading	Loads an undirected road network (unweighted in this case) from <code>roadNet-CA.txt</code> .	Matches the large-scale sparse graph setting discussed in the paper.
2.Initial SSSP	Performs standard Dijkstra from node 0 to all nodes.	Forms the <b>base shortest path tree (SPT)</b> .
3.Simulated Graph Updates	Applies edge deletions and insertions to simulate dynamic changes.	Matches the paper's edge update model: dynamic deletions and insertions.

<b>4.Affected Node Detection</b>	Identifies nodes whose shortest paths are impacted by the updates (i.e., whose parent is invalidated).	Matches the “Affected Vertex Identification” phase of the paper.
<b>5.Parallel Propagation</b>	Propagates new distances from affected nodes using OpenMP ( <code>#pragma omp parallel for</code> ).	This is the <b>Parallel Frontier Expansion</b> technique in the paper.
<b>6.Iteration Until Convergence</b>	Repeats propagation until no more distance changes occur ( <code>while (changed)</code> loop).	Aligns with the <b>template loop structure</b> in the paper: update → check → repeat.
<b>7.Speedup Evaluation</b>	Measures time taken for dynamic update and compares it to full recomputation using Dijkstra.	Demonstrates the benefit of the parallel update template: <b>reduced recomputation</b> .
<b>8. Metrics &amp; Logging</b>	Logs execution time, number of updated and unreachable nodes, and thread-wise performance to CSV.	Supports performance analysis as done in experimental sections of the paper.

## OpenMP-Specific Features Used

- `#pragma omp parallel for schedule(dynamic)`
  - Enables multithreaded processing of affected nodes during update phase.

- **#pragma omp critical**
  - Ensures that updates to shared `dist[]` and `parent[]` are thread-safe.

Our code **fully realizes the core algorithm** proposed in the research paper:

- Uses **multithreaded parallelism** on shared memory
- Implements **incremental updates** rather than recomputation
- Achieves **scalable and fast dynamic SSSP updates**
- Tracks and evaluates performance metrics (execution time, speedup, convergence)

The approach is especially suited for **real-time graph systems** where updates are frequent and full recomputation is impractical.

## OpenMP – Based Dynamic SSSP Implementation

### High-Level Goal

This code implements a **parallel SSSP solver using OpenCL**, based on a simplified Dijkstra-like method (relaxation over edges). It targets **GPU acceleration** and handles large graphs like `roadNet-CA`.

### Execution Flow

- ♦ **1. Graph Loading (Host - `main.cpp`)**
  - Loads an **unweighted undirected graph** from `roadNet-CA.txt`.

- Random weights are assigned between 1 and 100.
- Stores edges in three arrays: `edges_u`, `edges_v`, and `weights`.
- Initializes a `dist[ ]` vector with `INT_MAX`, except for the source node (0), which is set to 0.

**Why this matters:** Converts real-world road networks into a form usable by OpenCL kernels.

---

## ♦ 2. OpenCL Setup

- Initializes OpenCL environment:
  - Gets platform and GPU device.
  - Creates a context and command queue.
- Reads kernel code from `dijkstra.cl`.
- Compiles the kernel and checks for build errors.

**Why this matters:** Sets up the environment to run GPU-parallel code across thousands of edges.

---

## ♦ 3. Memory Management

- Allocates memory on the GPU using `clCreateBuffer()`:
  - Graph edge arrays (`edges_u`, `edges_v`, `weights`)
  - Distance array (`dist`)

- `updated` flag buffer (used to check convergence)
- Sets kernel arguments using `clSetKernelArg()`.

**Why this matters:** Transfers all graph and algorithm state to GPU memory.

---

#### ◆ 4. Kernel Execution Loop

- Launches the OpenCL kernel repeatedly **until no further distance updates occur**:
  1. Set `updated = 0`.
  2. Run kernel with one thread per edge.
  3. If `dist[]` is changed in any thread, kernel writes `updated = 1`.
  4. Repeat until `updated == 0`.

**Why this matters:** Implements relaxation rounds until no shorter paths are found.

---

#### ◆ 5. Final Output and Logging

- After convergence, the distance array is copied back from GPU to CPU.
- Measures and displays:
  - Execution time
  - Nodes reachable from the source
  - Distance to first 10 reachable nodes
  - Logs performance metrics to CSV.

## Kernel Logic (**dijkstra.cl**)

```
__kernel void dijkstra(...) {  
    int i = get_global_id(0); // Each thread handles edge i
```

1. Read edge ( $u \rightarrow v$ ) and weight.
2. If **dist[u]** is not **INT\_MAX**, compute **new\_dist = dist[u] + w**.
3. Atomically update **dist[v]** using **atomic\_min()** if a shorter path is found.
4. If **dist[v]** was changed, set **updated[0] = 1**.

**Why this matters:** Implements Dijkstra-like **edge relaxation** in parallel, using atomic operations for correctness.

---

## Comparison of OpenMP vs OpenCL vs METIS

Aspect	OpenMP (CPU)	OpenCL (GPU)	METIS (Partitioned CPU)
Platform	CPU (multi-threaded)	GPU (many-core)	CPU with METIS partitioning
Graph Model	Unweighted, undirected	Weighted (random), undirected	Weighted (random), undirected

<b>Algorithm Type</b>	Incremental Dijkstra update	Iterative parallel relaxation	Full Dijkstra over entire graph
<b>Parallelism Granularity</b>	Node-based (frontier expansion via OpenMP)	Edge-based (one thread per edge)	Per partition (parallel potential with post-processing)
<b>Affected Node Detection</b>	Yes (based on parent edge)	No (blind edge relaxations until convergence)	No (recomputes from scratch)
<b>Distance Updates</b>	SSSP tree repair with OpenMP	Atomic <code>min()</code> over edges per round	Full priority queue-based Dijkstra
<b>Termination Condition</b>	Converges when no affected nodes remain	Converges when <code>updated = 0</code> across kernel round	Runs once, no iteration
<b>Strengths</b>	Fastest for dynamic updates, thread-scalable	Leverages GPU massively, scalable to large graphs	Good for static graphs with prepartitioning
<b>Limitations</b>	Assumes shared memory; limited to one machine	Not exact Dijkstra (no priority queue), kernel setup overhead	High partitioning overhead, not dynamic
<b>Conforms to Paper?</b>	Fully matches (template: detect → propagate → fix)	Partially (parallel updates, no frontier detection)	Baseline only (full recomputation)



<b>Time Complexity</b>	$\sim O(k \cdot (V+E))$ localized updates	$\sim O(r \times E)$ where $r$ is number of rounds	$O(V \log V + E)$
<b>Best Use Case</b>	Frequent small dynamic updates	Massively parallel full SSSP	Static, large graphs in distributed environments

## Why Use METIS Over Manual Partitioning?

Using **METIS** for graph partitioning offers several technical advantages compared to manual (naive or round-robin) partitioning, especially in the context of **parallel graph processing** like SSSP.

## Real Benefits of METIS

1. **Fewer Inter-Partition Dependencies**  
→ Leads to **less synchronization** and **better parallelism** in distributed or threaded environments.
2. **Workload Balance**  
→ Ensures **each thread or node gets a similar amount of work**, preventing bottlenecks.
3. **Minimal Edge Cuts**  
→ Essential in SSSP and BFS where crossing partitions introduces coordination overhead.
4. **Ease of Use**  
→ METIS handles complex heuristics internally — no need to design custom logic.

## **When Manual Might Be Acceptable?**

- For small graphs with uniform structure
-